

# Deep Network Components & Training

# BASIC UNITS OF NEURAL NETWORK

**Input Layer:** Input variables, sometimes called the visible layer.

**Hidden Layers:** Layers of nodes between the input and output layers. There may be one or more of these layers.

**Output Layer:** A layer of nodes that produce the output variables.

Things to describe the shape and capability of a neural network:

- **Size:** The number of nodes in the model/network.
- **Width:** The number of nodes in a specific layer.
- **Depth:** The number of layers in a neural network (We don't count input as a layer).
- **Capacity:** The type or structure of functions that can be learned by a network configuration. Sometimes called "*representational capacity*".
- **Architecture:** The specific arrangement of the layers and nodes in the network.

# WHAT WE LEARN FIRST?

- Iteration, Batch, Epoch
- Recap of Activation Functions
- Data Preprocessing
- Weight Initialization
- Batch Normalization
- Hyperparameter Optimization

# EPOCH, BATCH, ITERATION

## Epoch



An Epoch represents one iteration over the entire dataset.

## Batch



We cannot pass the entire dataset into the neural network at once. So, we divide the dataset into number of batches.

## Iteration



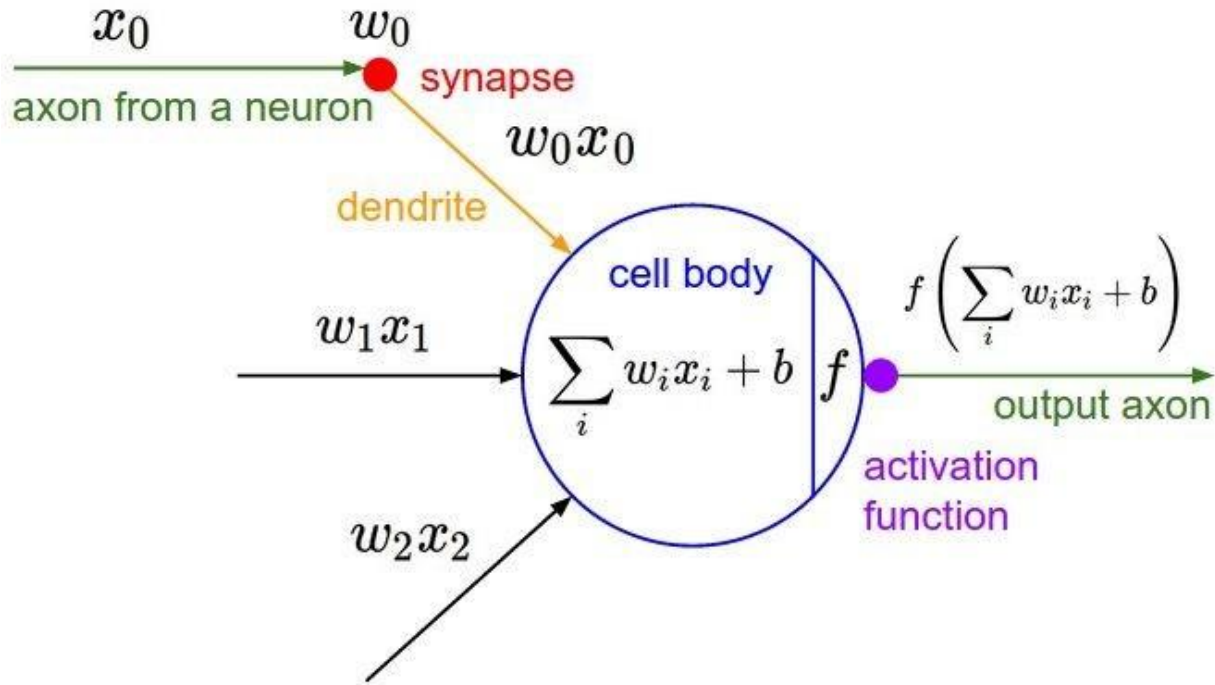
If we have 10,000 images as data and a batch size of 200, then an epoch should contain  $10,000/200 = 50$  iterations.

# EPOCH, BATCH, ITERATION



# Activation Functions

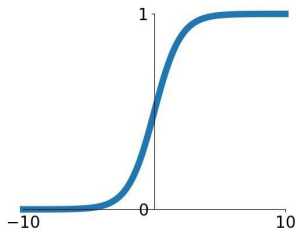
# ACTIVATION FUNCTIONS



# ACTIVATION FUNCTIONS

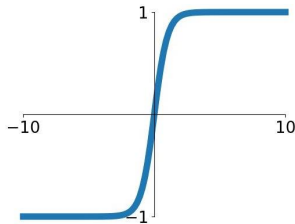
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



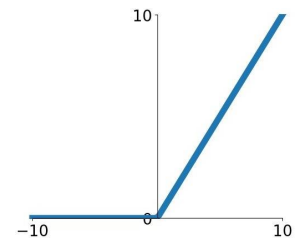
## tanh

$$\tanh(x)$$



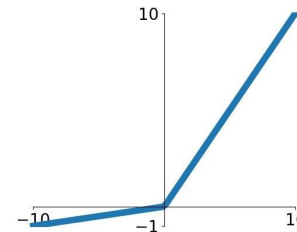
## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$

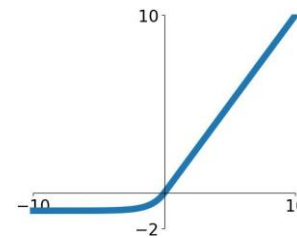


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

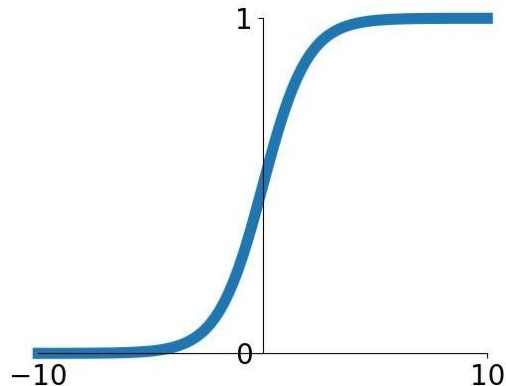
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$





# ACTIVATION FUNCTIONS

$$\sigma(x) = 1/(1 + e^{-x})$$



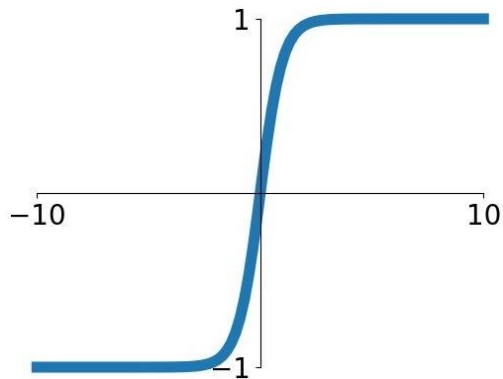
**Sigmoid**

- Squashes numbers to range [0,1]
- Historically popular since they have nice interpretation as a saturating “firing rate” of a neuron

3 problems:

1. Saturated neurons “kill” the gradients
2. Sigmoid outputs are not zero-centered
3.  $\exp()$  is a bit compute expensive

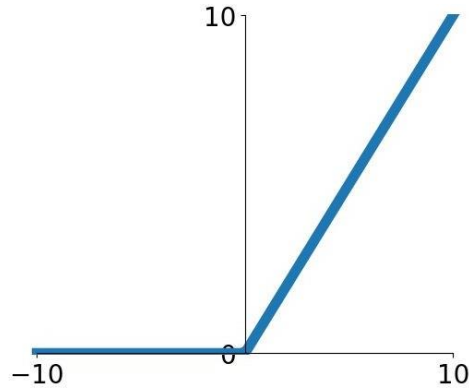
# ACTIVATION FUNCTIONS



$\tanh(x)$

- Squashes numbers to range  $[-1,1]$
- zero centered (nice)
- still kills gradients when saturated :(

# ACTIVATION FUNCTIONS



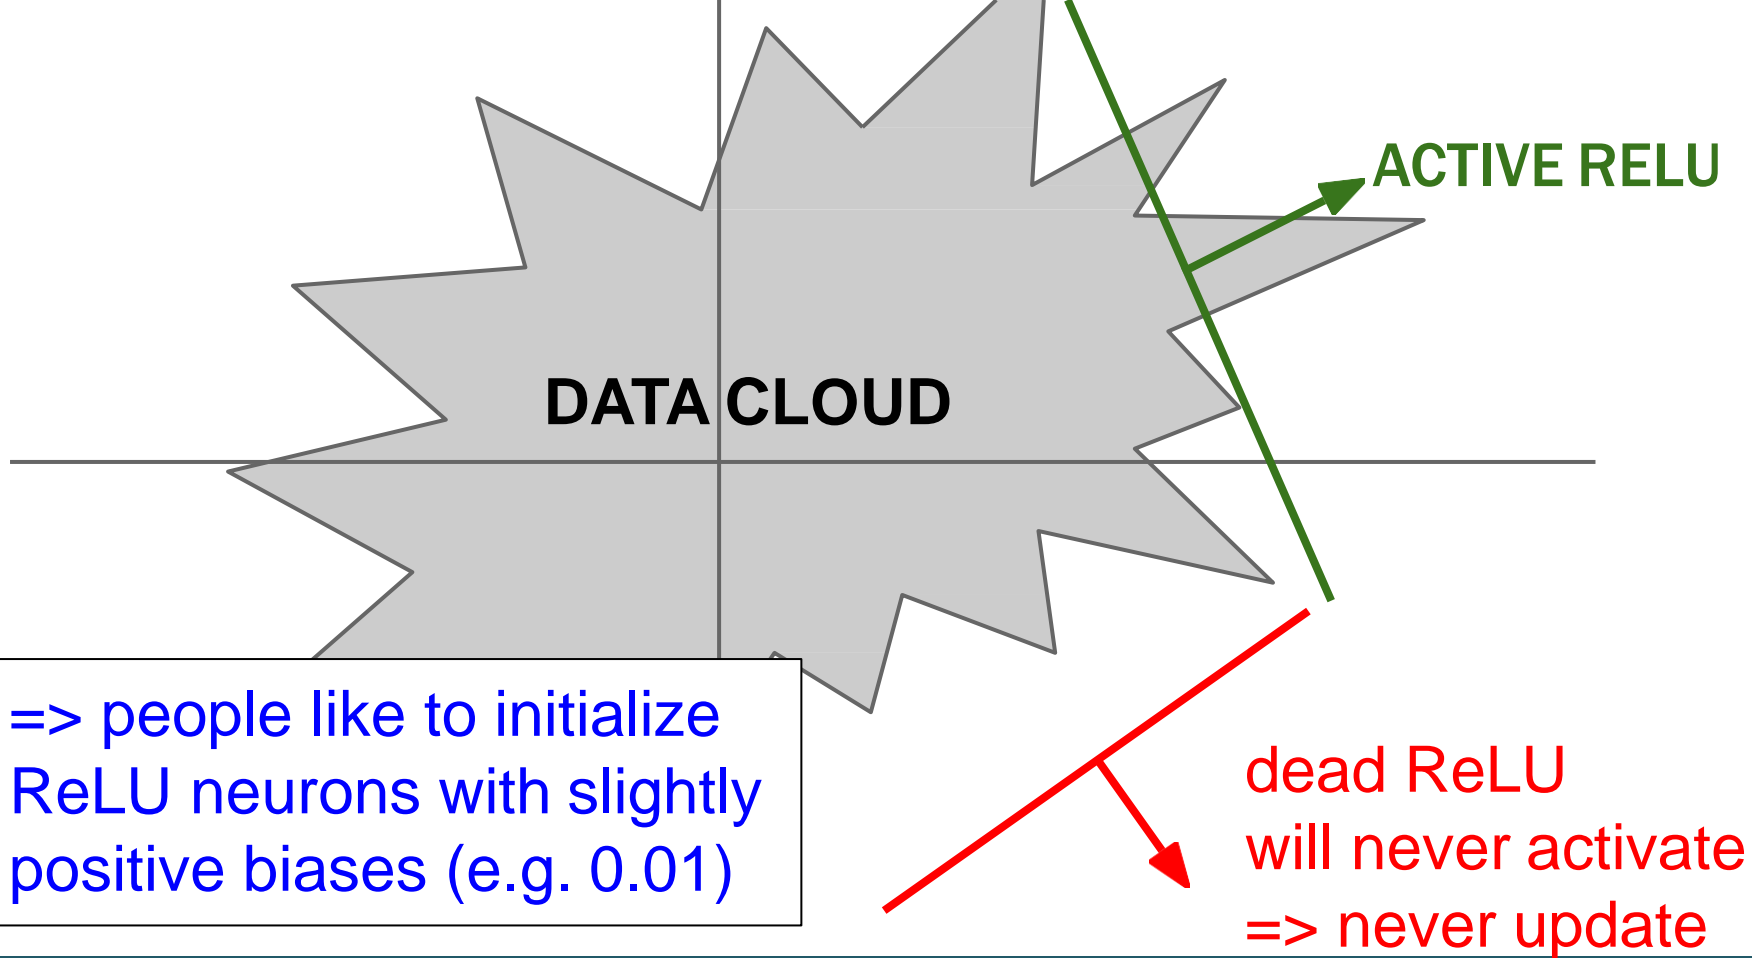
**ReLU**  
(Rectified Linear Unit)

- COMPUTES  $F(X) = \text{MAX}(0, X)$

- Does not saturate (in +region)
- Very computationally efficient
- Converges much faster than sigmoid/tanh in practice (e.g. 6x)
- Actually more biologically plausible than sigmoid

- Not zero-centered output
- An annoyance:

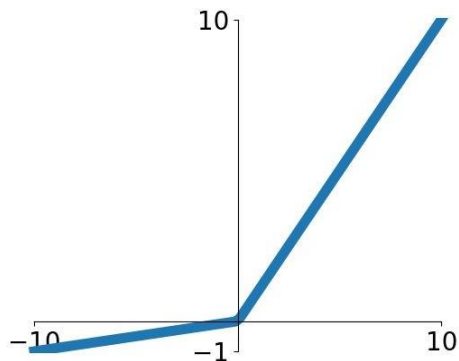
hint: what is the gradient when  $x < 0$ ?



# ACTIVATION FUNCTIONS

[Mass et al., 2013]

[He et al., 2015]



## Leaky ReLU

$$f(x) = \max(0.01x, x)$$

- Does not saturate
- Computationally efficient
- Converges much faster than sigmoid/tanh in practice! (e.g. 6x)
- **will not “die”.**

## Parametric Rectifier (PReLU)

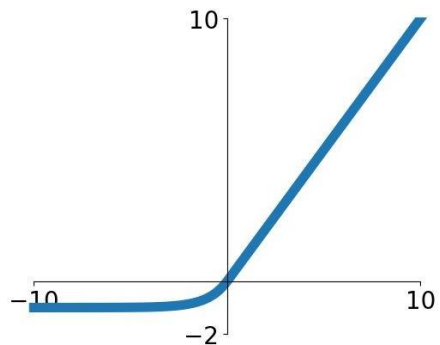
$$f(x) = \max(\alpha x, x)$$

backprop into  $\alpha$   
(parameter)

# ACTIVATION FUNCTIONS

[Clevert et al., 2015]

## Exponential Linear Units (ELU)



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$

- All benefits of ReLU
- Closer to zero mean outputs
- Negative saturation regime compared with Leaky ReLU adds some robustness to noise
- Computation requires  $\exp()$

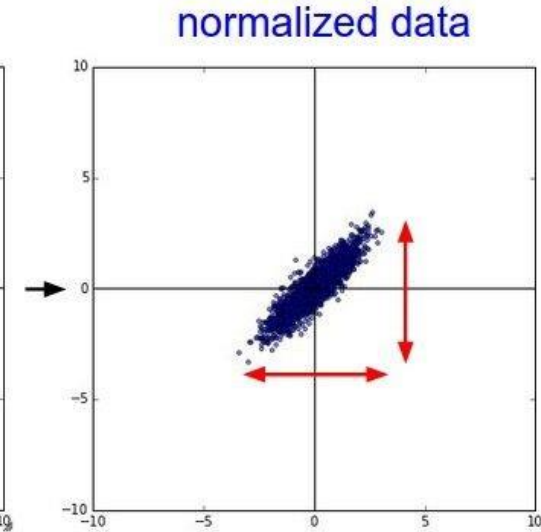
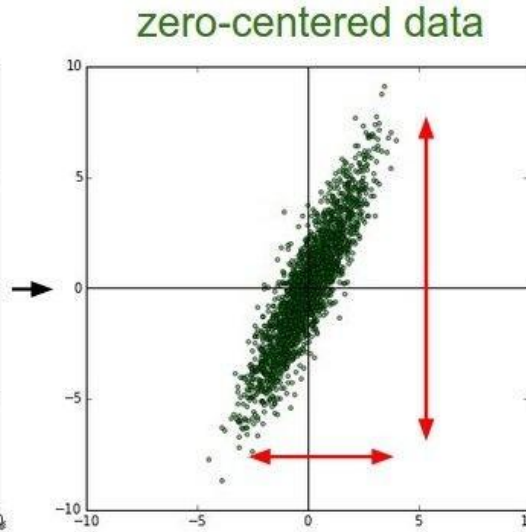
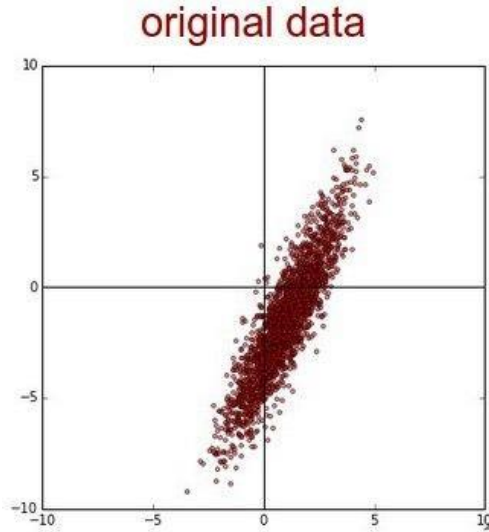
## IN PRACTICE:

- Use **ReLU**. Be careful with your learning rates
- Try out **Leaky ReLU / Maxout / ELU**
- Try out **tanh** but don't expect much
- **Don't use sigmoid**

# Data Preprocessing



# STEP 1: PREPROCESS THE DATA



```
X -= np.mean(X, axis = 0)
```

```
X /= np.std(X, axis = 0)
```

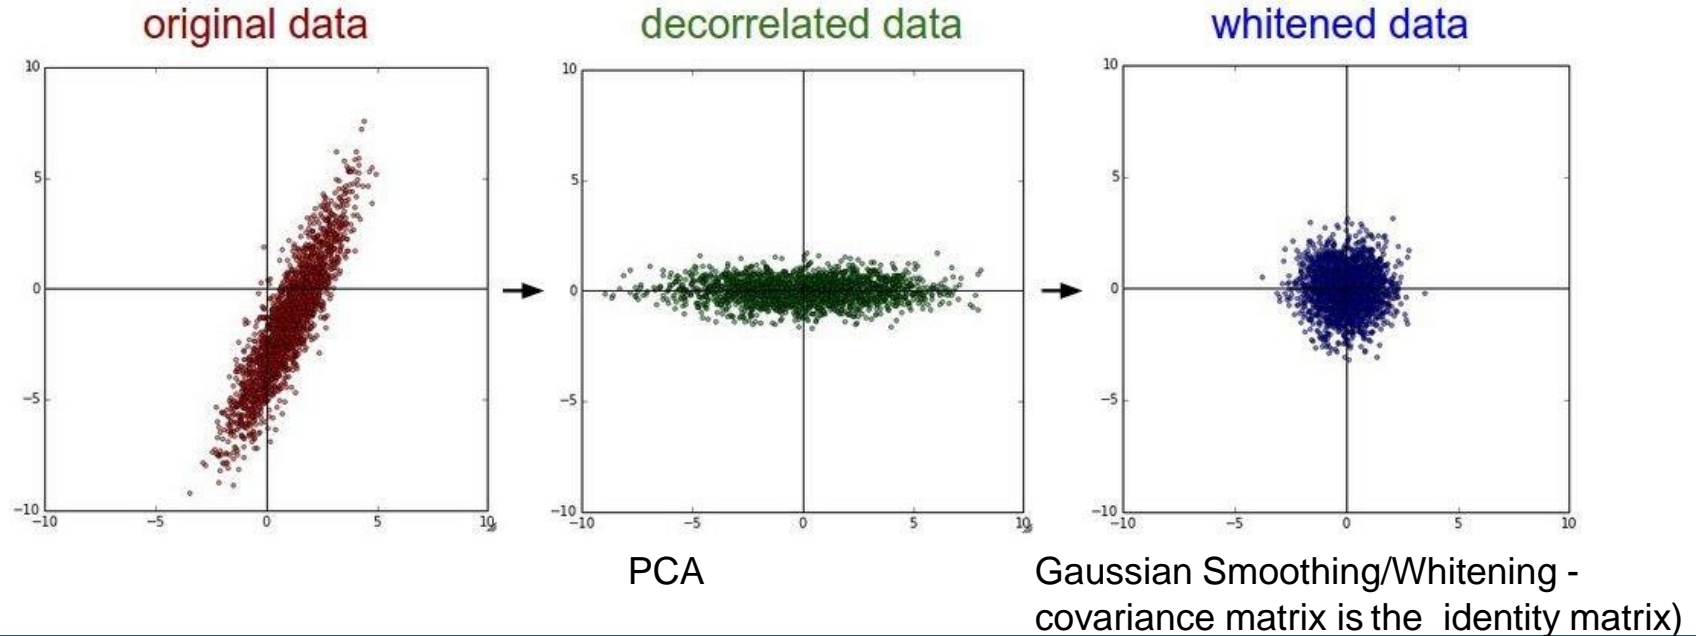
(Assume  $X$  [NxD] is data matrix, each example in a row)

Subtract every sample with mean

Then, divide every sample With standard deviation

# STEP 1: PREPROCESS THE DATA

In practice, you may also see **PCA** and **Whitening** of the data



# IN PRACTICE FOR IMAGES: CENTER ONLY

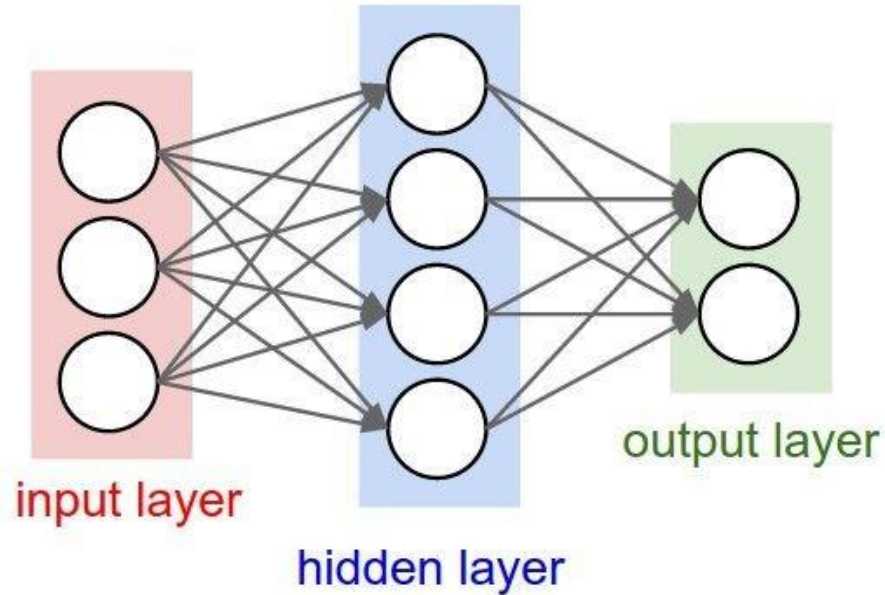
e.g. consider CIFAR-10 example with  $[32,32,3]$  images

- Subtract the mean image (e.g. AlexNet)  
(mean image =  $[32,32,3]$  array)
- Subtract per-channel mean (e.g. VGGNet)  
(mean along each channel = 3 numbers)

Not common to normalize  
variance, to do PCA

# Weight Initialization

- Q: WHAT HAPPENS WHEN  $W=0$  INIT IS USED? NO LEARNING



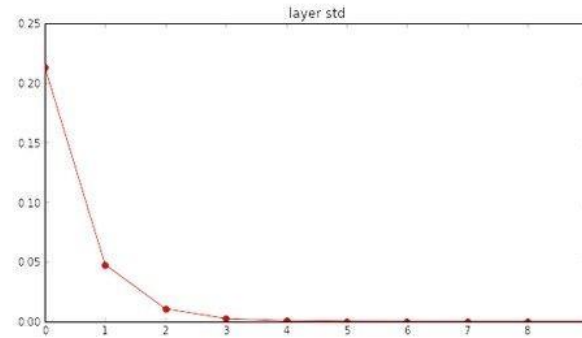
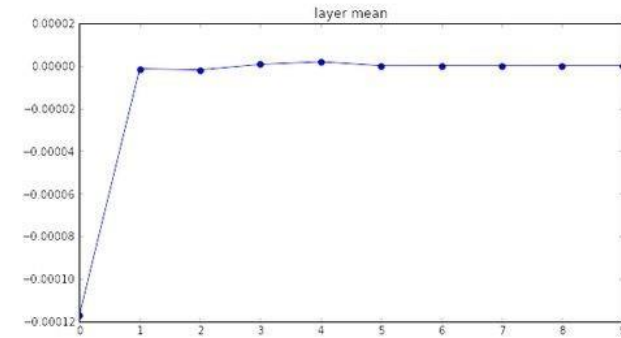
- First idea: **Small random numbers**  
(gaussian with zero mean and  $1e-2$  standard deviation)

```
W = 0.01* np.random.randn(D,H)
```

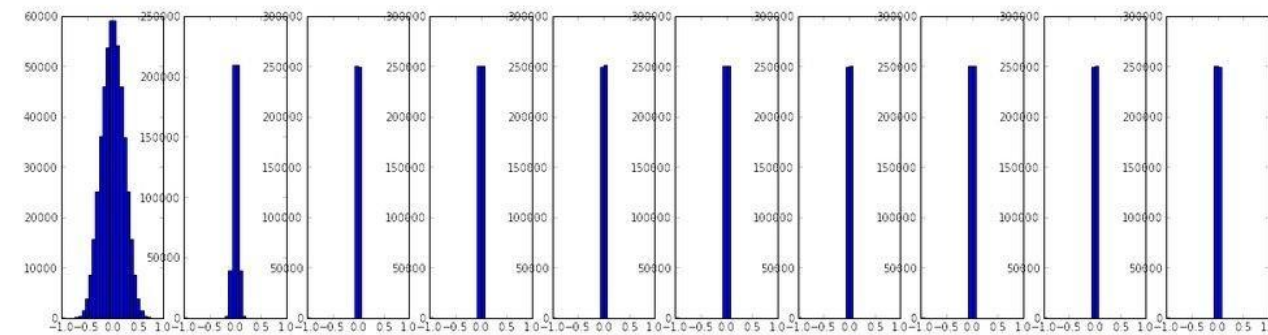
Works ~okay for small networks, but problems with deeper networks.

Input layer had mean 0.000927 and std 0.998388  
 hidden layer 1 had mean -0.000117 and std 0.213081  
 hidden layer 2 had mean -0.000001 and std 0.047551  
 hidden layer 3 had mean -0.000002 and std 0.010630  
 hidden layer 4 had mean 0.000001 and std 0.002378  
 hidden layer 5 had mean 0.000002 and std 0.000532  
 hidden layer 6 had mean -0.000000 and std 0.000119  
 hidden layer 7 had mean 0.000000 and std 0.000026  
 hidden layer 8 had mean -0.000000 and std 0.000006  
 hidden layer 9 had mean 0.000000 and std 0.000001  
 hidden layer 10 had mean -0.000000 and std 0.000000

**ALL ACTIVATIONS  
BECOME ZERO!**



Q: think about the backward pass.  
What do the gradients look like?

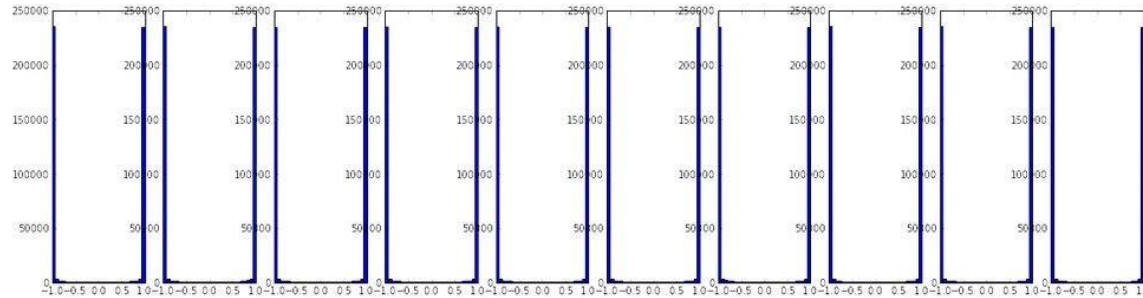
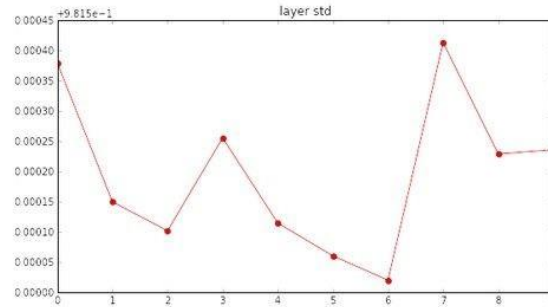
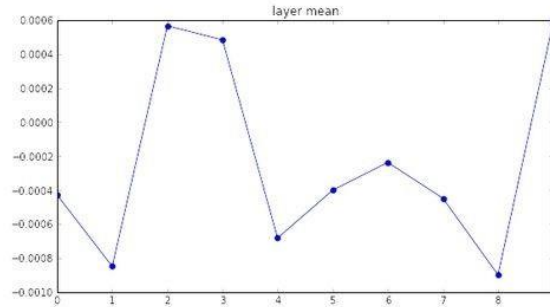


Hint: think about backward pass for a  $W \cdot X$  gate.

```
W = np.random.randn(fan_in, fan_out) * 1.0 # layer initialization
```

input layer had mean 0.001800 and std 1.001311  
hidden layer 1 had mean -0.000430 and std 0.981879  
hidden layer 2 had mean -0.000849 and std 0.981649  
hidden layer 3 had mean 0.000566 and std 0.981601  
hidden layer 4 had mean 0.000483 and std 0.981755  
hidden layer 5 had mean -0.000682 and std 0.981614  
hidden layer 6 had mean -0.000401 and std 0.981560  
hidden layer 7 had mean -0.000237 and std 0.981520  
hidden layer 8 had mean -0.000448 and std 0.981913  
hidden layer 9 had mean -0.000899 and std 0.981728  
hidden layer 10 had mean 0.000584 and std 0.981736

\*1.0 instead of \*0.01



**ALMOST ALL NEURONS  
COMPLETELY  
SATURATED, EITHER -1  
AND 1. GRADIENTS  
WILL BE ALL ZERO.**



# XAVIOR INITILAIZATION

Xavier initialization sets a layer's weights to values chosen from a random uniform distribution that's bounded between:

$$\pm \frac{\sqrt{6}}{\sqrt{n_i + n_{i+1}}}$$

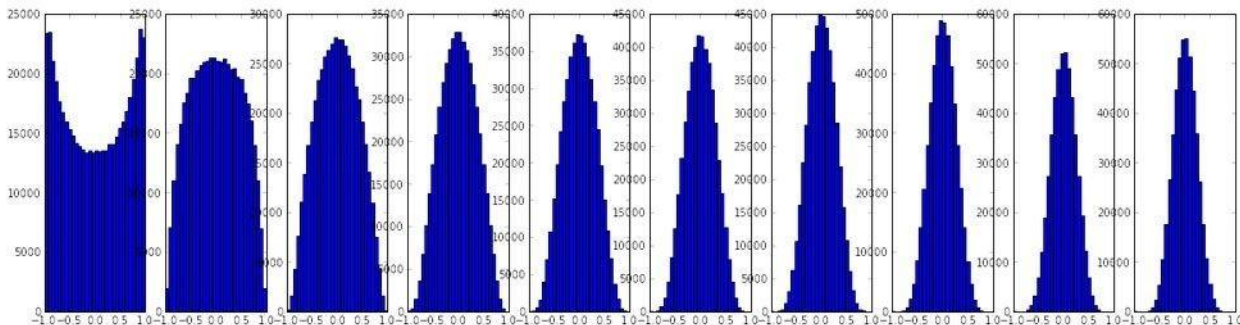
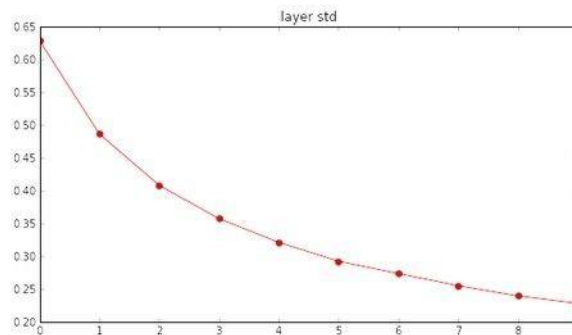
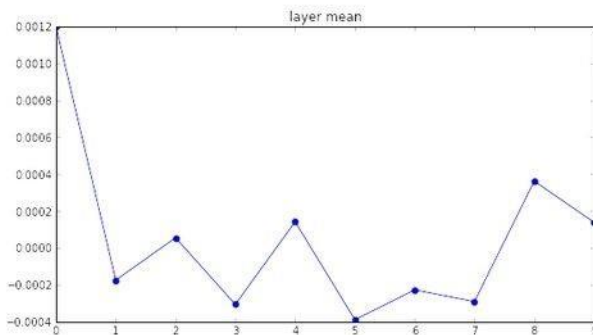
Where,

- $n_i$  is the number of incoming network connections, or “fan-in,” to the layer, and
- $n_{i+1}$  is the number of outgoing network connections from that layer, also known as the “fan-out.”

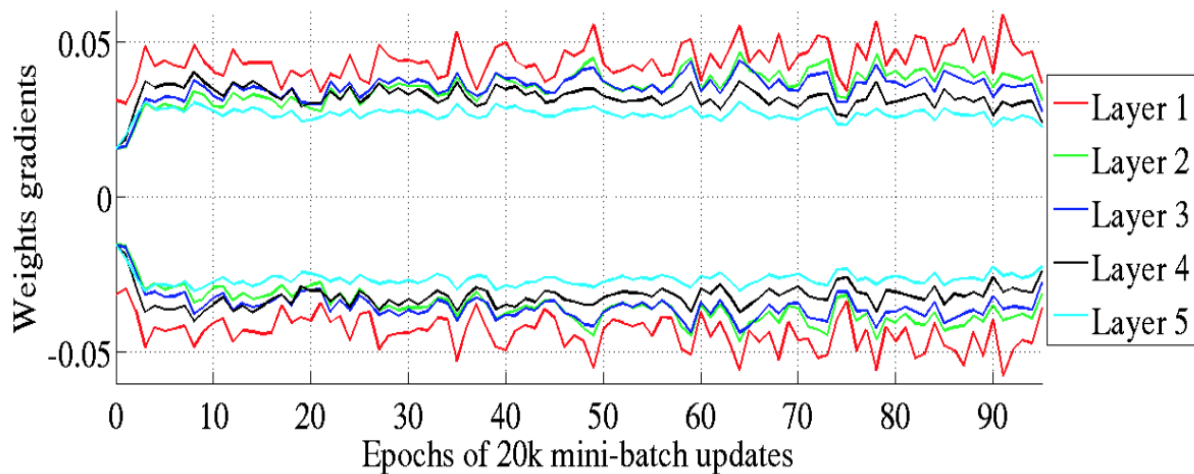
# “XAVIER INITIALIZATION” [GLOROT ET AL., 2010]

Reasonable initialization.

input layer had mean 0.001800 and std 1.001311  
hidden layer 1 had mean 0.001198 and std 0.627953  
hidden layer 2 had mean -0.000175 and std 0.486051  
hidden layer 3 had mean 0.000055 and std 0.407723  
hidden layer 4 had mean -0.000306 and std 0.357108  
hidden layer 5 had mean 0.000142 and std 0.320917  
hidden layer 6 had mean -0.000389 and std 0.292116  
hidden layer 7 had mean -0.000228 and std 0.273387  
hidden layer 8 had mean -0.000291 and std 0.254935  
hidden layer 9 had mean 0.000361 and std 0.239266  
hidden layer 10 had mean 0.000139 and std 0.228008

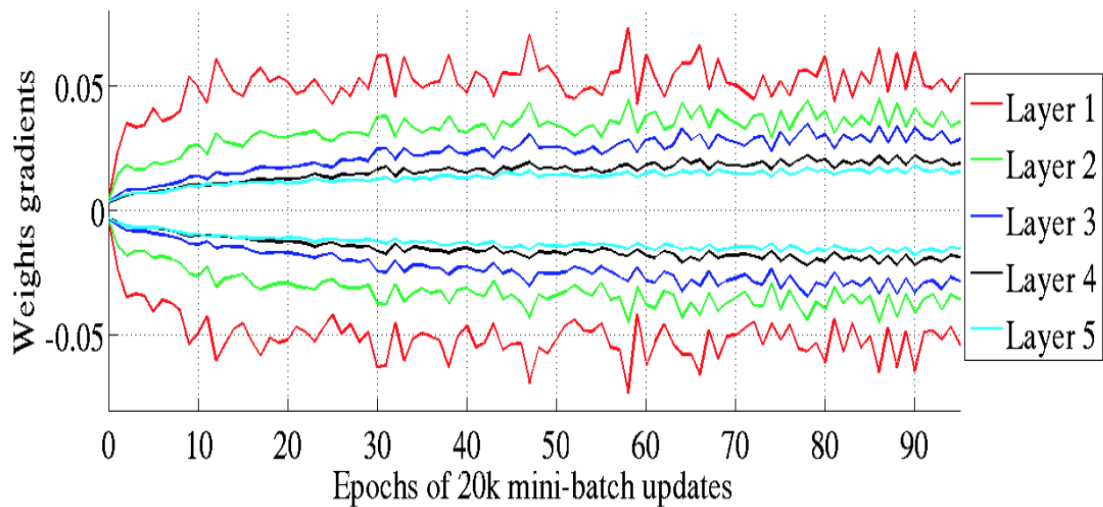


Glorot and Bengio believed that Xavier weight initialization would maintain the variance of activations and back-propagated gradients all the way up or down the layers of a network. In their experiments they observed that Xavier initialization enabled a 5-layer network to maintain near identical variances of its weight gradients across layers.



With Xavier init. Credit: [Glorot & Bengio](#).

Conversely, it turned out that using “standard” initialization brought about a much bigger gap in variance between weight gradients at the network’s lower layers, which were higher, and those at its top-most layers, which were approaching zero.



Without Xavier init. Credit: [Glorot & Bengio](#).

# PROPER INITIALIZATION IS AN ACTIVE AREA OF RESEARCH...

***Understanding the difficulty of training deep feedforward neural networks***

by Glorot and Bengio, 2010

***Exact solutions to the nonlinear dynamics of learning in deep linear neural networks*** by

Saxe et al, 2013

***Random walk initialization for training very deep feedforward networks*** by Sussillo and

Abbott, 2014

***Delving deep into rectifiers: Surpassing human-level performance on ImageNet***

***classification*** by He et al., 2015

***Data-dependent Initializations of Convolutional Neural Networks*** by Krähenbühl et al., 2015

***All you need is a good init***, Mishkin and Matas, 2015

...

# Batch Normalization

# BATCH NORMALIZATION IS NOT NORMALIZATION

- Normalization is a data preparation technique to make sure that the data is consistent and has a uniform scale. In Images, commonly used data preparation techniques include pixel normalization and image normalization.
- For example, when we have features from 0 to 1 and some from 1 to 1000, we should normalize them to speed up learning.
- If the input layer is benefiting from it, why not do the same thing also for the values in the hidden layers, that are changing all the time, and get 10 times or more improvement in the training speed.

# BATCH NORMALIZATION – WHY?

- One of the important issues with using neural network is that the **training of the network takes a long time** for effectively deep networks.

## Backpropagation:

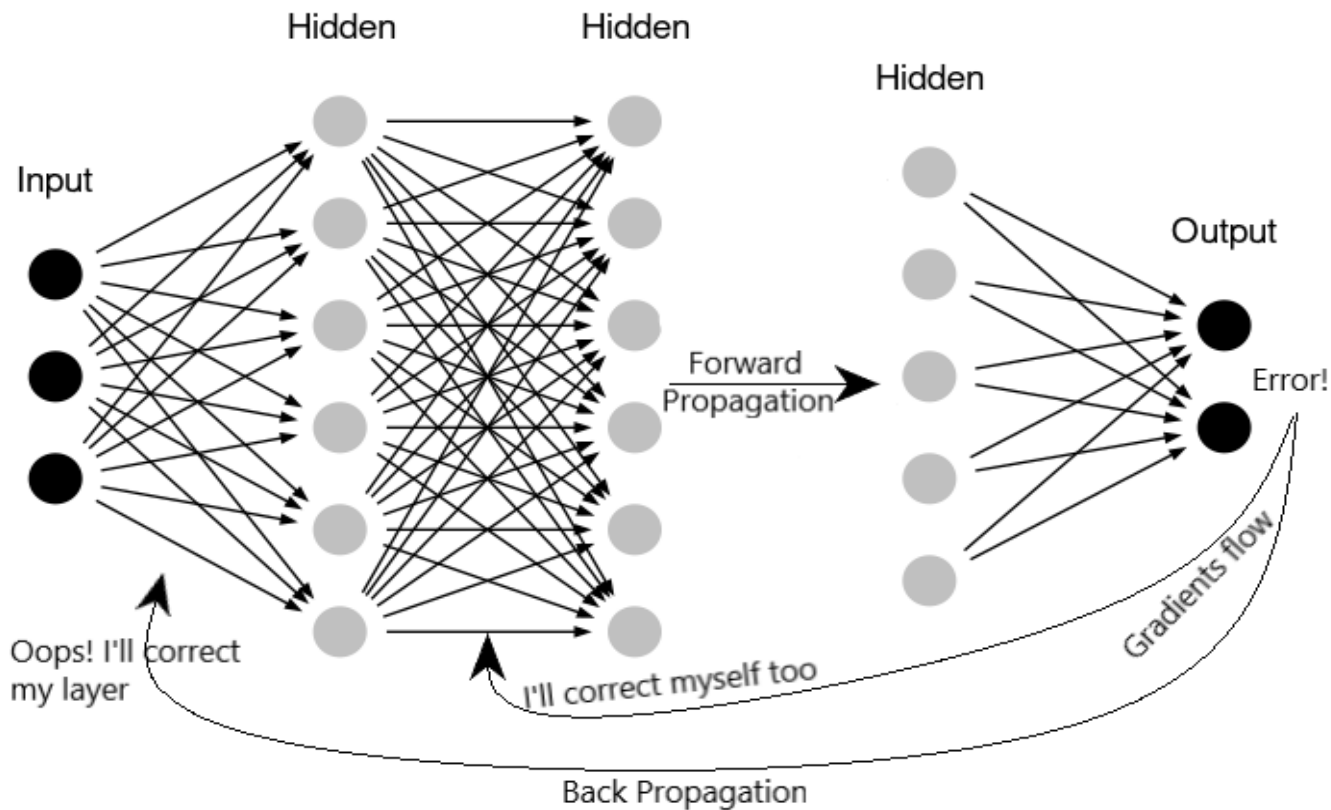
- Neural networks learn the problem using BackPropagation algorithm.
- By backpropagation, the neurons learn how much error they did and correct themselves, i.e, correct their “weights” and “biases”.
- By this they learn the problem to produce correct outputs given the inputs.
- BackPropagation involves computing gradients for each layer and propagating it backward, hence the name.

## Problem in Backpropagation:

During backpropagation of errors to the weights and biases, we'll face a undesired property of **Internal Covariate Shift**. This makes the network too long to train.

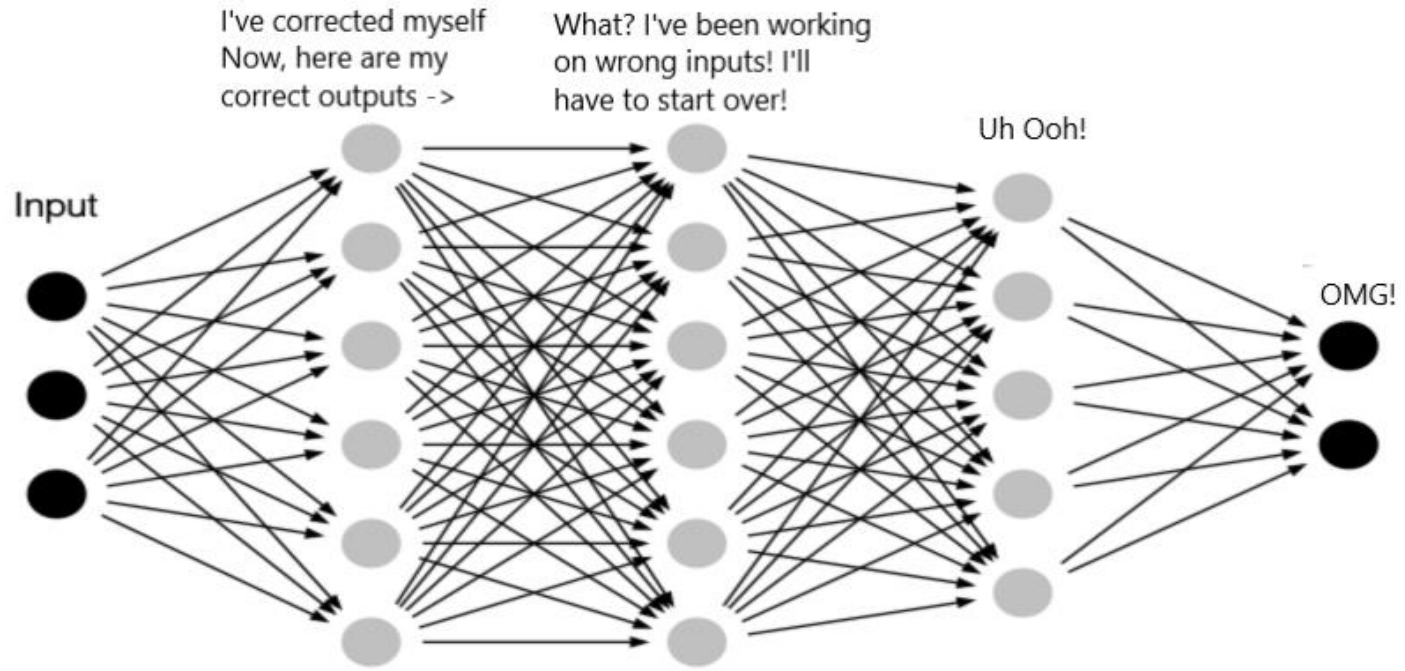


# BATCH NORMALIZATION – WHY?



# BATCH NORMALIZATION – TO OVERCOME INTERNAL COVARIANCE SHIFT

For example, in the network given below, the 2nd layer adjusts its weights and biases to correct for the output. But due to this readjustment, the output of 2nd layer, i.e, the input of 3rd layer is changed for same initial input. So the third layer has to learn from scratch to produce the correct outputs for the same data.

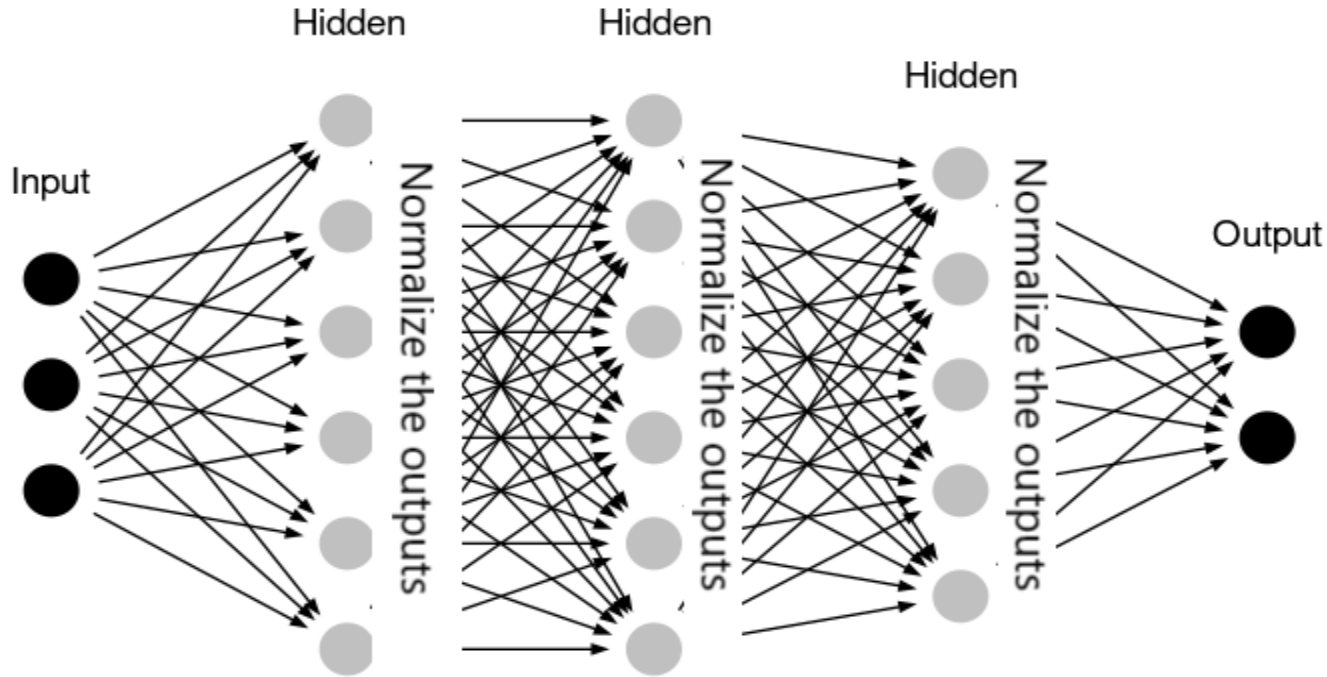


# BATCH NORMALIZATION – INTERNAL COVARIANCE SHIFT

- This presents the problem of a layer starting to learn after it's previous layer, i.e, 3rd layer learns after 2nd finished, 4th starts learning after 3rd, etc.
- Similarly, think of the current existing deep neural networks that are about 100 to even 1000 layers deep! It would really take epochs to train them.
- More specifically, due to changes in weights of previous layers, the distribution of input values for current layer changes, forcing it to learn from new “input distribution”.

# BATCH NORMALIZATION – HOW IT WORKS?

- Batch Normalization adds Normalization “layer” between each layers.
- An important thing to note here is that **normalization has to be done separately for each dimension** (input neuron), **over the ‘mini-batches’**, and not altogether with all dimensions. **Hence the name ‘batch’ normalization**.



# BATCH NORMALIZATION – HOW IT WORKS?

Due to this normalization “layers” between each fully connected layers, the range of input distribution of each layer stays the same, no matter the changes in the previous layer. Given  $x$  inputs from  $k$ -th neuron.

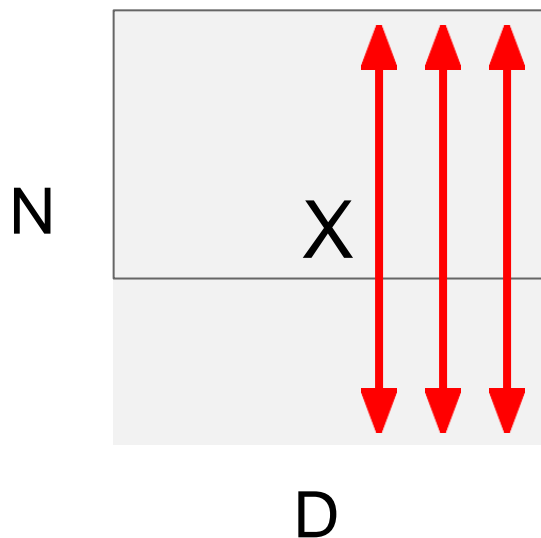
$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

Normalization brings all the inputs centered around 0. This way, there is not much change in each layer input. So, layers in the network can learn from the back-propagation simultaneously, without waiting for the previous layer to learn. This fastens up the training of networks.

# BATCH NORMALIZATION

[Ioffe and Szegedy, 2015]

“you want unit gaussian activations?  
just make them so.”



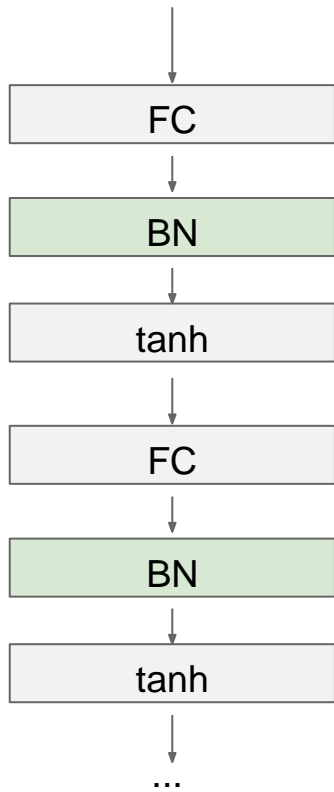
1. compute the empirical mean and variance independently for each dimension.

2. Normalize (divide by standard deviation)

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

# BATCH NORMALIZATION

[Ioffe and Szegedy, 2015]



Usually inserted after Fully Connected or Convolutional layers, and before nonlinearity.

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

# BATCH NORMALIZATION

[Ioffe and Szegedy, 2015]

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_{1\dots m}\}$ ;  
Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

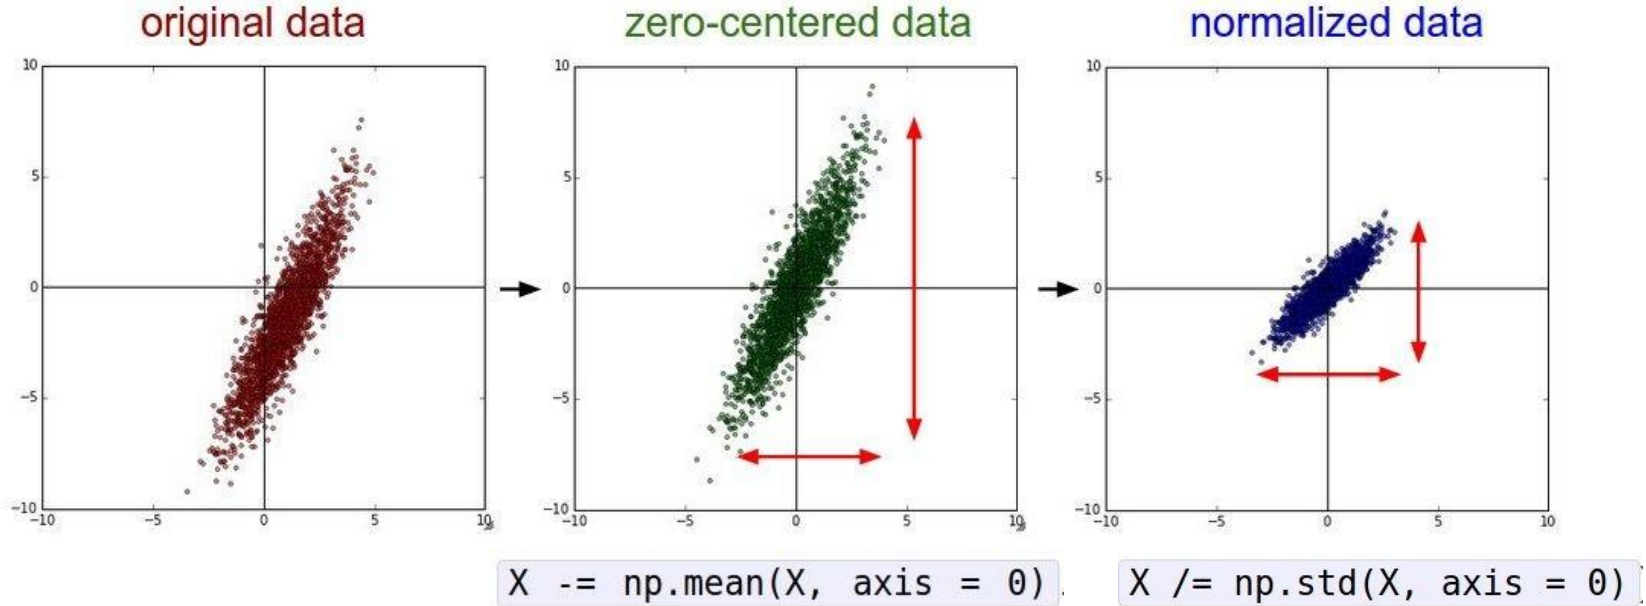
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

- Improves gradient flow through the network
- Allows higher learning rates
- Reduces the strong dependence on initialization
- Acts as a form of regularization in a funny way, and slightly reduces the need for dropout, maybe



# Babysitting the Learning Process

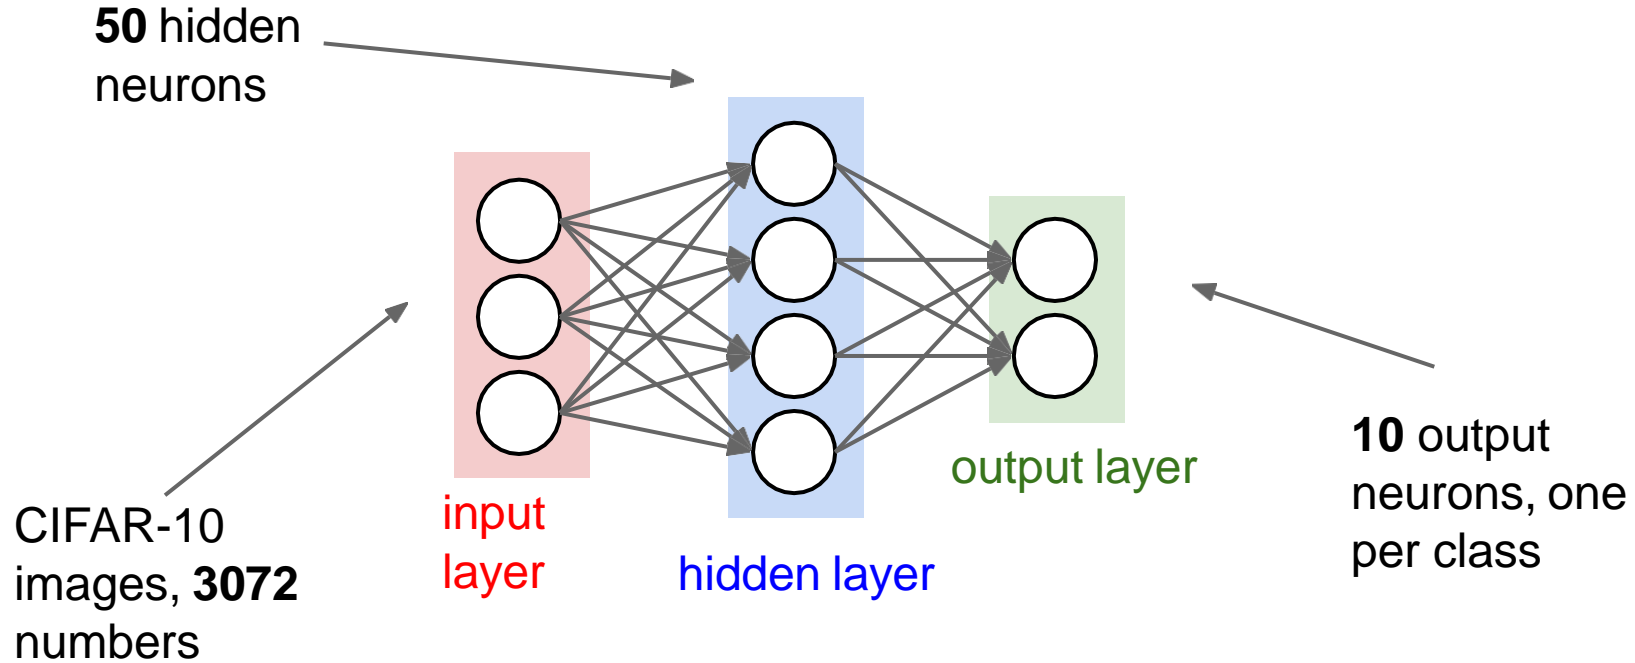
# STEP 1: PREPROCESS THE DATA



(Assume  $X$  [NxD] is data matrix,  
each example in a row)

## STEP 2: CHOOSE THE ARCHITECTURE:

SAY WE START WITH ONE HIDDEN LAYER OF 50 NEURONS:



## DOUBLE CHECK THAT THE LOSS IS REASONABLE:

```
def init_two_layer_model(input_size, hidden_size, output_size):  
    # initialize a model  
    model = {}  
    model['W1'] = 0.0001 * np.random.randn(input_size, hidden_size)  
    model['b1'] = np.zeros(hidden_size)  
    model['W2'] = 0.0001 * np.random.randn(hidden_size, output_size)  
    model['b2'] = np.zeros(output_size)  
    return model
```

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes  
loss, grad = two_layer_net(X_train, model, y_train, 0.0)  
print loss
```

disable regularization

2.30261216167

loss ~2.3.  
"correct" for  
10 classes

returns the loss and the  
gradient for all parameters

## DOUBLE CHECK THAT THE LOSS IS REASONABLE:

```
def init_two_layer_model(input_size, hidden_size, output_size):  
    # initialize a model  
    model = {}  
    model['W1'] = 0.0001 * np.random.randn(input_size, hidden_size)  
    model['b1'] = np.zeros(hidden_size)  
    model['W2'] = 0.0001 * np.random.randn(hidden_size, output_size)  
    model['b2'] = np.zeros(output_size)  
    return model
```

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes  
loss, grad = two_layer_net(X_train, model, y_train, 1e3) crank up regularization  
print loss
```

3.06859716482

loss went up, good. (sanity check)

# Lets try to train now...

**Tip:** Make sure that you can overfit very small portion of the training data

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
X_tiny = X_train[:20] # take 20 examples
y_tiny = y_train[:20]
best_model, stats = trainer.train(X_tiny, y_tiny, X_tiny, y_tiny,
                                  model, two_layer_net,
                                  num_epochs=200, reg=0.0,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = False,
                                  learning_rate=1e-3, verbose=True)
```

The above code:

- take the first 20 examples from CIFAR-10
- turn off regularization (reg = 0.0)
- use simple vanilla 'sgd'

# Lets try to train now...

**Tip:** Make sure that you can overfit very small portion of the training data

Very small loss,  
train accuracy 1.00,  
nice!

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
X_tiny = X_train[:20] # take 20 examples
y_tiny = y_train[:20]
best_model, stats = trainer.train(X_tiny, y_tiny, X_tiny, y_tiny,
                                  model, two_layer_net,
                                  num_epochs=200, reg=0.0,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = False,
                                  learning_rate=1e-3, verbose=True)
```

```
Finished epoch 1 / 200: cost 2.302603, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 2 / 200: cost 2.302258, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 3 / 200: cost 2.301849, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 4 / 200: cost 2.301196, train: 0.650000, val 0.650000, lr 1.000000e-03
Finished epoch 5 / 200: cost 2.300044, train: 0.650000, val 0.650000, lr 1.000000e-03
Finished epoch 6 / 200: cost 2.297864, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 7 / 200: cost 2.293595, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 8 / 200: cost 2.285096, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 9 / 200: cost 2.268094, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 10 / 200: cost 2.234787, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 11 / 200: cost 2.173187, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 12 / 200: cost 2.076862, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 13 / 200: cost 1.974090, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 14 / 200: cost 1.895885, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 15 / 200: cost 1.820876, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 16 / 200: cost 1.737430, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 17 / 200: cost 1.642356, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 18 / 200: cost 1.535239, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 19 / 200: cost 1.421527, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 20 / 200: cost 1.305760, train: 0.650000, val 0.650000, lr 1.000000e-03
```

```
Finished epoch 195 / 200: cost 0.002694, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 196 / 200: cost 0.002674, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 197 / 200: cost 0.002655, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 198 / 200: cost 0.002635, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 199 / 200: cost 0.002617, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 200 / 200: cost 0.002597, train: 1.000000, val 1.000000, lr 1.000000e-03
finished optimization. best validation accuracy: 1.000000
```



## Lets try to train now...

Start with small  
regularization and find  
learning rate that  
makes the loss go  
down.

[illegible]



# LETS TRY TO TRAIN NOW...

Start with small regularization and find learning rate that makes the loss go down.

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e-6, verbose=True)
```

Epoch	Cost	Train Loss	Val Loss	Learning Rate
1 / 10	2.302576	0.080000	0.103000	1.000000e-06
2 / 10	2.302582	0.121000	0.124000	1.000000e-06
3 / 10	2.302558	0.119000	0.138000	1.000000e-06
4 / 10	2.302519	0.127000	0.151000	1.000000e-06
5 / 10	2.302517	0.158000	0.171000	1.000000e-06
6 / 10	2.302518	0.179000	0.172000	1.000000e-06
7 / 10	2.302466	0.180000	0.176000	1.000000e-06
8 / 10	2.302452	0.175000	0.185000	1.000000e-06
9 / 10	2.302459	0.206000	0.192000	1.000000e-06
10 / 10	2.302420	0.190000	0.192000	1.000000e-06

finished optimization. best validation accuracy: 0.192000

Loss barely changing

# LETS TRY TO TRAIN NOW...

Start with small  
regularization and find  
learning rate that  
makes the loss go  
down.

**loss not going down:**  
learning rate too low

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e-6, verbose=True)
```

Epoch	Cost	Train Loss	Val Loss	Learning Rate
1 / 10	2.302576	0.080000	0.103000	1.000000e-06
2 / 10	2.302582	0.121000	0.124000	1.000000e-06
3 / 10	2.302558	0.119000	0.138000	1.000000e-06
4 / 10	2.302519	0.127000	0.151000	1.000000e-06
5 / 10	2.302517	0.158000	0.171000	1.000000e-06
6 / 10	2.302518	0.179000	0.172000	1.000000e-06
7 / 10	2.302466	0.180000	0.176000	1.000000e-06
8 / 10	2.302452	0.175000	0.185000	1.000000e-06
9 / 10	2.302459	0.206000	0.192000	1.000000e-06
10 / 10	2.302420	0.190000	0.192000	1.000000e-06

finished optimization. best validation accuracy: 0.192000

Loss barely changing: Learning rate is  
probably too low

# LETS TRY TO TRAIN NOW...

Start with small regularization and find learning rate that makes the loss go down.

**loss not going down:**  
learning rate too low

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e-6, verbose=True)
```

Epoch	Cost	Train Accuracy	Val Accuracy	Learning Rate
1 / 10	2.302576	0.080000	0.103000	1.000000e-06
2 / 10	2.302582	0.121000	0.124000	1.000000e-06
3 / 10	2.302558	0.119000	0.138000	1.000000e-06
4 / 10	2.302519	0.127000	0.151000	1.000000e-06
5 / 10	2.302517	0.158000	0.171000	1.000000e-06
6 / 10	2.302518	0.179000	0.172000	1.000000e-06
7 / 10	2.302466	0.180000	0.176000	1.000000e-06
8 / 10	2.302452	0.175000	0.185000	1.000000e-06
9 / 10	2.302459	0.206000	0.192000	1.000000e-06
10 / 10	2.302420	0.190000	0.192000	1.000000e-06

finished optimization. best validation accuracy: 0.192000

Loss barely changing: Learning rate is probably too low

Notice train/val accuracy goes to 20% though, what's up with that? (remember this is softmax)

# LETS TRY TO TRAIN NOW...

Start with small  
regularization a  
find learning ra  
that makes the  
go down.

loss not going  
down:

learning rate too low

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
```

Now let's try learning rate 1e6.



Lets try to train now...

Start with small  
regularization and find  
learning rate that  
makes the loss go  
down.

**loss not going down:**  
learning rate too low  
**loss exploding:**  
learning rate too high

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e6, verbose=True)
```

```
/home/karpathy/cs231n/code/cs231n/classifiers/neural_net.py:50: RuntimeWarning: divide by zero en
countered in log
    data_loss = -np.sum(np.log(probs[range(N), y])) / N
/home/karpathy/cs231n/code/cs231n/classifiers/neural_net.py:48: RuntimeWarning: invalid value enc
ountered in subtract
    probs = np.exp(scores - np.max(scores, axis=1, keepdims=True))
Finished epoch 1 / 10: cost nan, train: 0.091000, val 0.087000, lr 1.000000e+06
Finished epoch 2 / 10: cost nan, train: 0.095000, val 0.087000, lr 1.000000e+06
Finished epoch 3 / 10: cost nan, train: 0.100000, val 0.087000, lr 1.000000e+06
```

cost: NaN almost  
always means high  
learning rate...



# LETS TRY TO TRAIN NOW...

Start with small regularization and find learning rate that makes the loss go down.

**loss not going down:**  
learning rate too low  
**loss exploding:**  
learning rate too high

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=3e-3, verbose=True)
```

```
Finished epoch 1 / 10: cost 2.186654, train: 0.308000, val 0.306000, lr 3.000000e-03
Finished epoch 2 / 10: cost 2.176230, train: 0.330000, val 0.350000, lr 3.000000e-03
Finished epoch 3 / 10: cost 1.942257, train: 0.376000, val 0.352000, lr 3.000000e-03
Finished epoch 4 / 10: cost 1.827868, train: 0.329000, val 0.310000, lr 3.000000e-03
Finished epoch 5 / 10: cost inf, train: 0.128000, val 0.128000, lr 3.000000e-03
Finished epoch 6 / 10: cost inf, train: 0.144000, val 0.147000, lr 3.000000e-03
```

3e-3 is still too high. Cost explodes....

=> Rough range for learning rate we should be cross-validating is somewhere [1e-3 ... 1e-5]

# Hyperparameter Optimization

# CROSS-VALIDATION STRATEGY

**coarse -> fine** cross-validation in stages

**First stage:** only a few epochs to get rough idea of what params work

**Second stage:** longer running time, finer search

... (repeat as necessary)

Tip for detecting explosions in the solver:

If the cost is ever  $> 3 \times$  original cost, break out early



# FOR EXAMPLE: RUN COARSE SEARCH

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)

    trainer = ClassifierTrainer()
    model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
    trainer = ClassifierTrainer()
    best_model_local, stats = trainer.train(X_train, y_train, X_val, y_val,
                                           model, two_layer_net,
                                           num_epochs=5, reg=reg,
                                           update='momentum', learning_rate_decay=0.9,
                                           sample_batches = True, batch_size = 100,
                                           learning_rate=lr, verbose=False)
```

← in log space!

```
val_acc: 0.412000, lr: 1.405206e-04, reg: 4.793564e-01, (1 / 100)
val_acc: 0.214000, lr: 7.231888e-06, reg: 2.321281e-04, (2 / 100)
val_acc: 0.208000, lr: 2.119571e-06, reg: 8.011857e+01, (3 / 100)
val_acc: 0.196000, lr: 1.551131e-05, reg: 4.374936e-05, (4 / 100)
val_acc: 0.079000, lr: 1.753300e-05, reg: 1.200424e+03, (5 / 100)
val_acc: 0.223000, lr: 4.215128e-05, reg: 4.196174e+01, (6 / 100)
val_acc: 0.441000, lr: 1.750259e-04, reg: 2.110807e-04, (7 / 100)
val_acc: 0.241000, lr: 6.749231e-05, reg: 4.226413e+01, (8 / 100)
val_acc: 0.482000, lr: 4.296863e-04, reg: 6.642555e-01, (9 / 100)
val_acc: 0.079000, lr: 5.401602e-06, reg: 1.599828e+04, (10 / 100)
val_acc: 0.154000, lr: 1.618508e-06, reg: 4.925252e-01, (11 / 100)
```

nice →

# NOW RUN FINER SEARCH...

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)
```

adjust range



```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-4, 0)
    lr = 10**uniform(-3, -4)
```

val_acc: 0.527000,	lr: 5.340517e-04,	reg: 4.097824e-01,	(0 / 100)
val_acc: 0.492000,	lr: 2.279484e-04,	reg: 9.991345e-04,	(1 / 100)
val_acc: 0.512000,	lr: 8.680827e-04,	reg: 1.349727e-02,	(2 / 100)
val_acc: 0.461000,	lr: 1.028377e-04,	reg: 1.220193e-02,	(3 / 100)
val_acc: 0.460000,	lr: 1.113730e-04,	reg: 5.244309e-02,	(4 / 100)
val_acc: 0.498000,	lr: 9.477776e-04,	reg: 2.001293e-03,	(5 / 100)
val_acc: 0.469000,	lr: 1.484369e-04,	reg: 4.328313e-01,	(6 / 100)
val_acc: 0.522000,	lr: 5.586261e-04,	reg: 2.312685e-04,	(7 / 100)
val_acc: 0.530000,	lr: 5.808183e-04,	reg: 8.259964e-02,	(8 / 100)
val_acc: 0.489000,	lr: 1.979168e-04,	reg: 1.010889e-04,	(9 / 100)
val_acc: 0.490000,	lr: 2.036031e-04,	reg: 2.406271e-03,	(10 / 100)
val_acc: 0.475000,	lr: 2.021162e-04,	reg: 2.287807e-01,	(11 / 100)
val_acc: 0.460000,	lr: 1.135527e-04,	reg: 3.905040e-02,	(12 / 100)
val_acc: 0.515000,	lr: 6.947668e-04,	reg: 1.562808e-02,	(13 / 100)
val_acc: 0.531000,	lr: 9.471549e-04,	reg: 1.433895e-03,	(14 / 100)
val_acc: 0.509000,	lr: 3.140888e-04,	reg: 2.857518e-01,	(15 / 100)
val_acc: 0.514000,	lr: 6.438349e-04,	reg: 3.033781e-01,	(16 / 100)
val_acc: 0.502000,	lr: 3.921784e-04,	reg: 2.707126e-04,	(17 / 100)
val_acc: 0.509000,	lr: 9.752279e-04,	reg: 2.850865e-03,	(18 / 100)
val_acc: 0.500000,	lr: 2.412048e-04,	reg: 4.997821e-04,	(19 / 100)
val_acc: 0.466000,	lr: 1.319314e-04,	reg: 1.189915e-02,	(20 / 100)
val_acc: 0.516000,	lr: 8.039527e-04,	reg: 1.528291e-02,	(21 / 100)

**53%** - relatively good  
for a 2-layer neural net  
with 50 hidden neurons.

# NOW RUN FINER SEARCH...

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)
```

adjust range

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-4, 0)
    lr = 10**uniform(-3, -4)
```

val_acc: 0.527000, lr: 5.340517e-04, reg: 4.097824e-01, (0 / 100)
val_acc: 0.492000, lr: 2.279484e-04, reg: 9.991345e-04, (1 / 100)
val_acc: 0.512000, lr: 8.680827e-04, reg: 1.349727e-02, (2 / 100)
val_acc: 0.461000, lr: 1.028377e-04, reg: 1.220193e-02, (3 / 100)
val_acc: 0.460000, lr: 1.113730e-04, reg: 5.244309e-02, (4 / 100)
val_acc: 0.498000, lr: 9.477776e-04, reg: 2.001293e-03, (5 / 100)
val_acc: 0.469000, lr: 1.484369e-04, reg: 4.328313e-01, (6 / 100)
val_acc: 0.522000, lr: 5.586261e-04, reg: 2.312685e-04, (7 / 100)
val_acc: 0.530000, lr: 5.808183e-04, reg: 8.259964e-02, (8 / 100)
val_acc: 0.489000, lr: 1.979168e-04, reg: 1.010889e-04, (9 / 100)
val_acc: 0.490000, lr: 2.036031e-04, reg: 2.406271e-03, (10 / 100)
val_acc: 0.475000, lr: 2.021162e-04, reg: 2.287807e-01, (11 / 100)
val_acc: 0.460000, lr: 1.135527e-04, reg: 3.905040e-02, (12 / 100)
val_acc: 0.515000, lr: 6.947668e-04, reg: 1.562808e-02, (13 / 100)
val_acc: 0.531000, lr: 9.471549e-04, reg: 1.433895e-03, (14 / 100)
val_acc: 0.509000, lr: 3.140888e-04, reg: 2.857518e-01, (15 / 100)
val_acc: 0.514000, lr: 6.438349e-04, reg: 3.033781e-01, (16 / 100)
val_acc: 0.502000, lr: 3.921784e-04, reg: 2.707126e-04, (17 / 100)
val_acc: 0.509000, lr: 9.752279e-04, reg: 2.850865e-03, (18 / 100)
val_acc: 0.500000, lr: 2.412048e-04, reg: 4.997821e-04, (19 / 100)
val_acc: 0.466000, lr: 1.319314e-04, reg: 1.189915e-02, (20 / 100)
val_acc: 0.516000, lr: 8.039527e-04, reg: 1.528291e-02, (21 / 100)

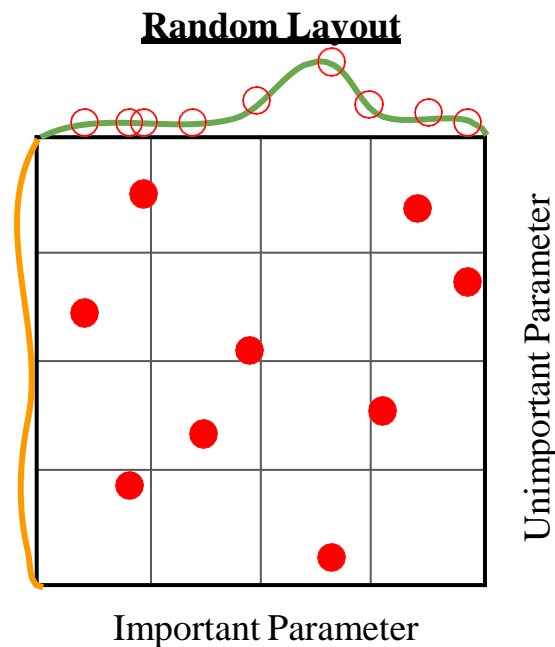
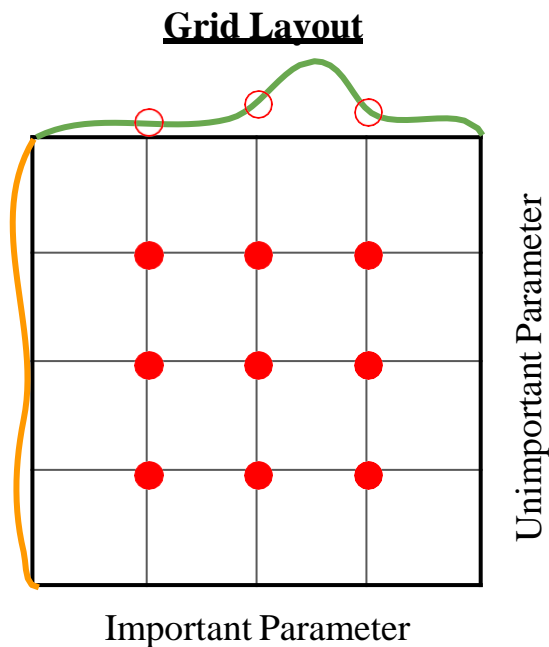
**53%** - relatively good  
for a 2-layer neural net  
with 50 hidden neurons.

But this best  
cross-validation result is  
worrying. Why?



# Random Search vs. Grid Search

*RANDOM  
SEARCH FOR  
Hyper-Parameter Optimization  
Bergstra and Bengio, 2012*





# HYPERPARAMETERS TO PLAY WITH:

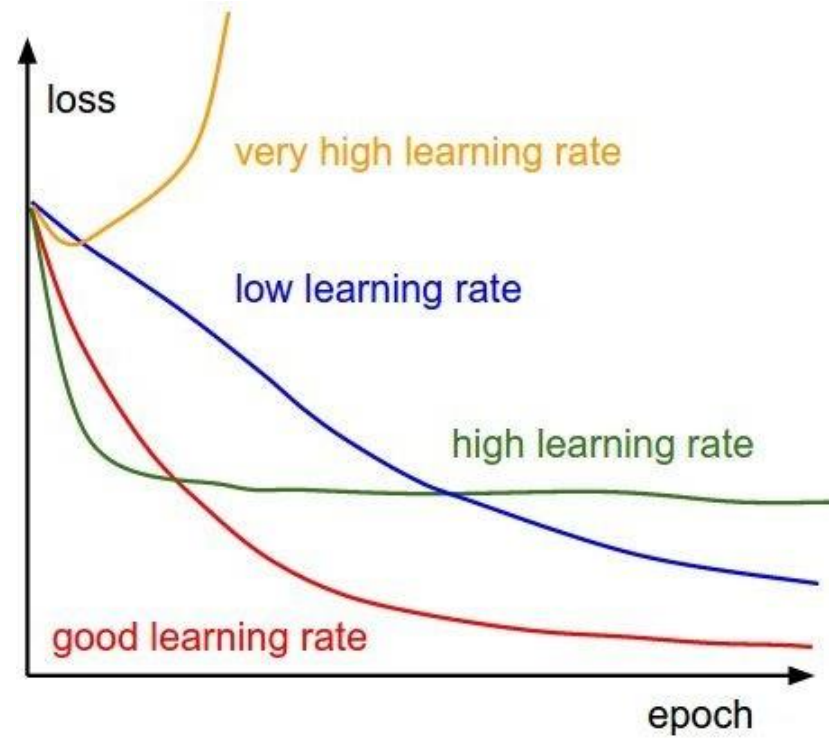
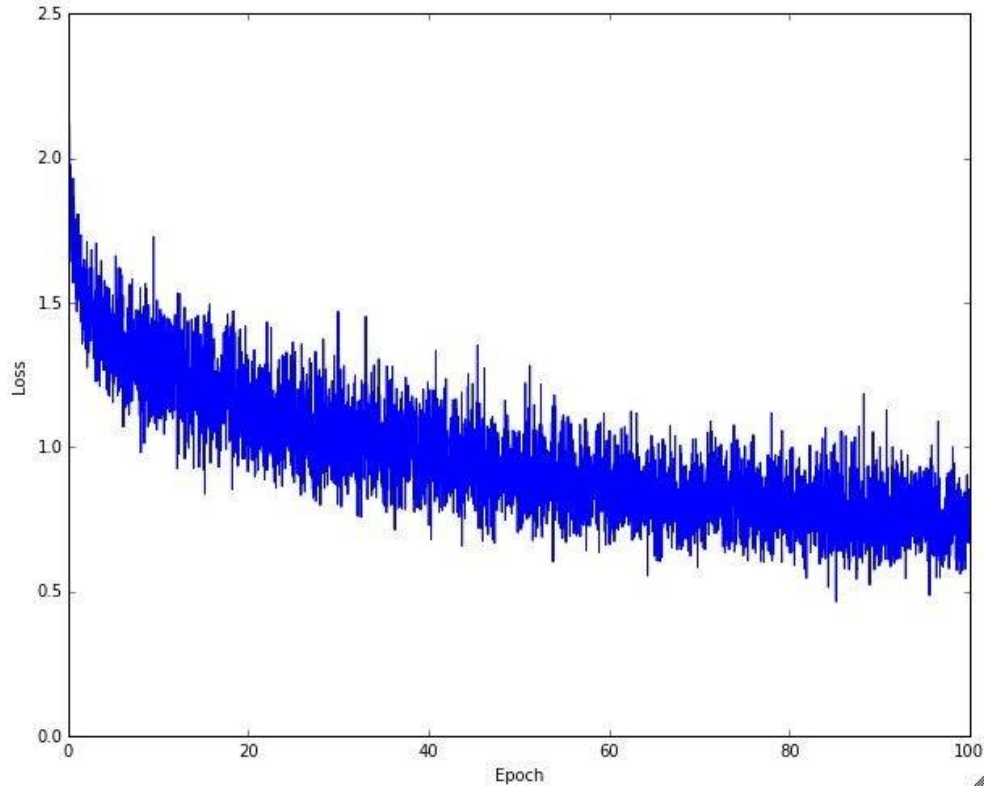
- Network architecture
- learning rate, its decay schedule, update type
- regularization (L2/Dropout strength)

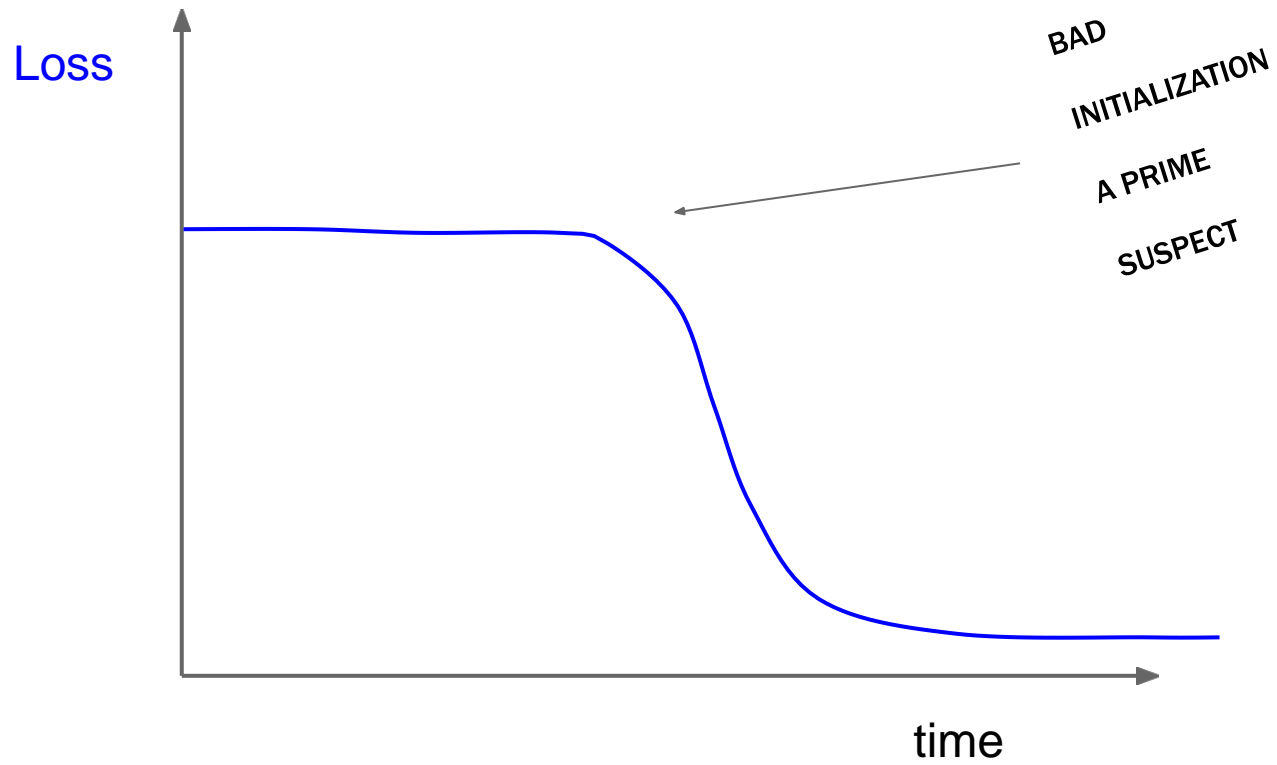
neural networks practitioner  
music = loss function



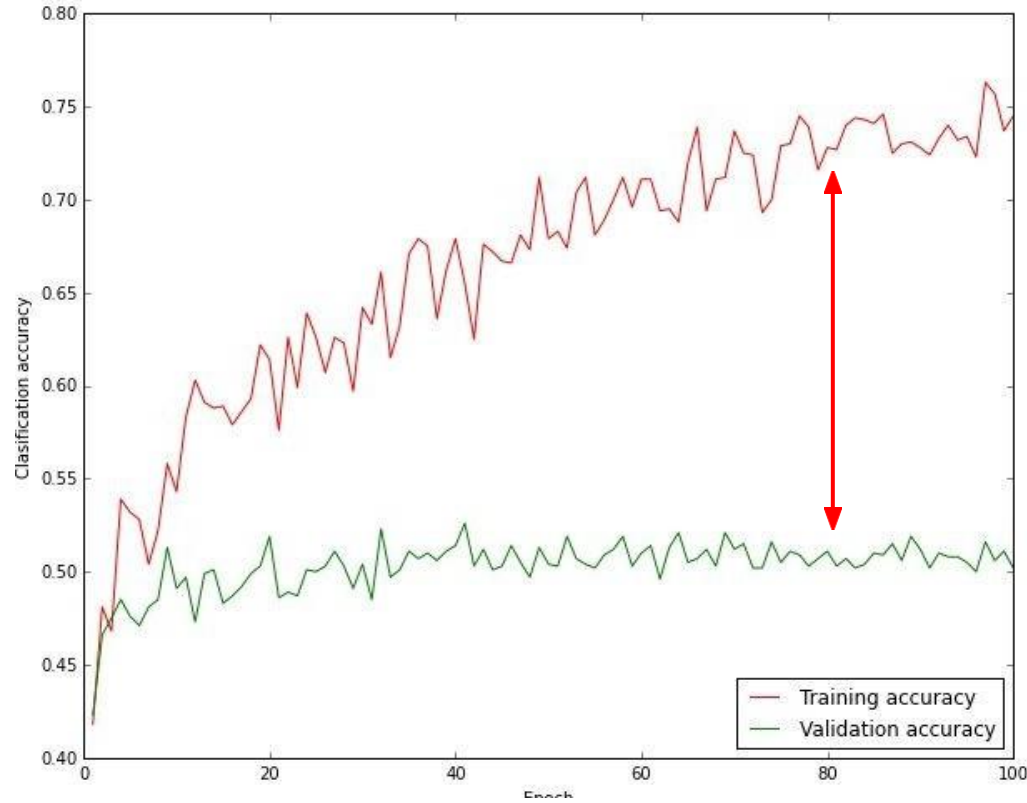
[This image](#) by Paolo Guereta is licensed under [CC-BY 2.0](#)

# MONITOR AND VISUALIZE THE LOSS CURVE





# MONITOR AND VISUALIZE THE ACCURACY:



big gap = overfitting

=> increase regularization strength?

no gap

=> increase model capacity?



# SUMMARY

We looked in detail at:

- Activation Functions (use ReLU)
- Data Preprocessing (images: subtract mean)
- Weight Initialization (use Xavier init)
- Batch Normalization (use)
- Babysitting the Learning process
- Hyperparameter Optimization  
(random sample hyperparams, in log space when appropriate)

## NOW, LET US UNDERSTAND THE TRAINING PROCESS IN DEEP...FURTHER DETAILS

- Parameter update schemes/Optimization (Will be covered in the coming classes)
- Regularization (Dropout etc.)
- Evaluation (Cross Validation, Ensembles etc.)

# Ensembles

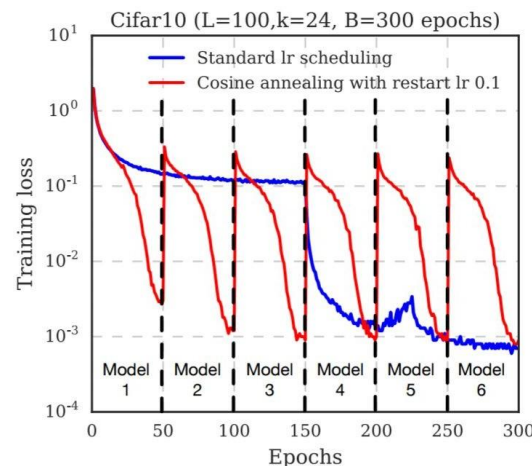
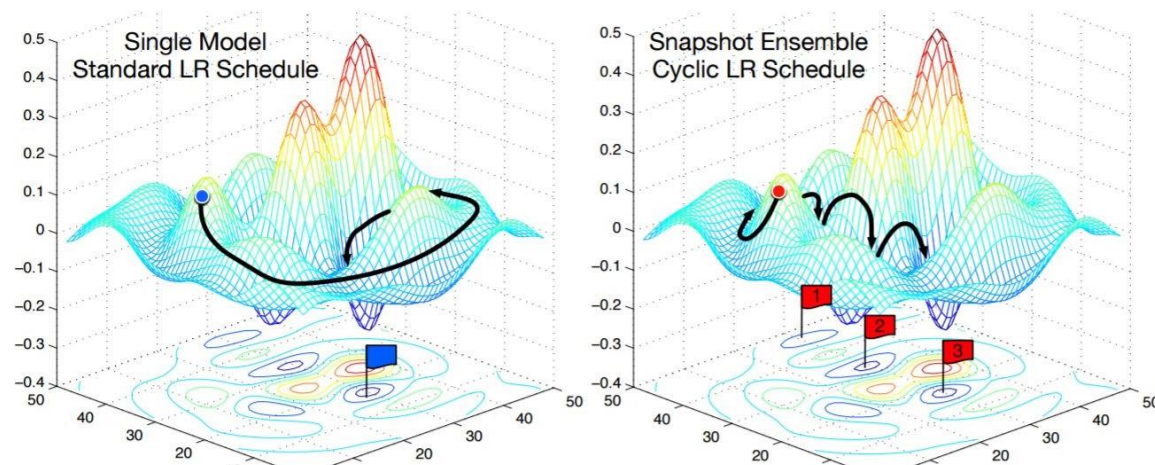
# MODEL ENSEMBLES

1. Train multiple independent models
2. At test time average their results  
(Take average of predicted probability distributions, then choose argmax)

Enjoy 2% extra performance

# MODEL ENSEMBLES: TIPS AND TRICKS

Instead of training independent models, use multiple snapshots of a single model during training!



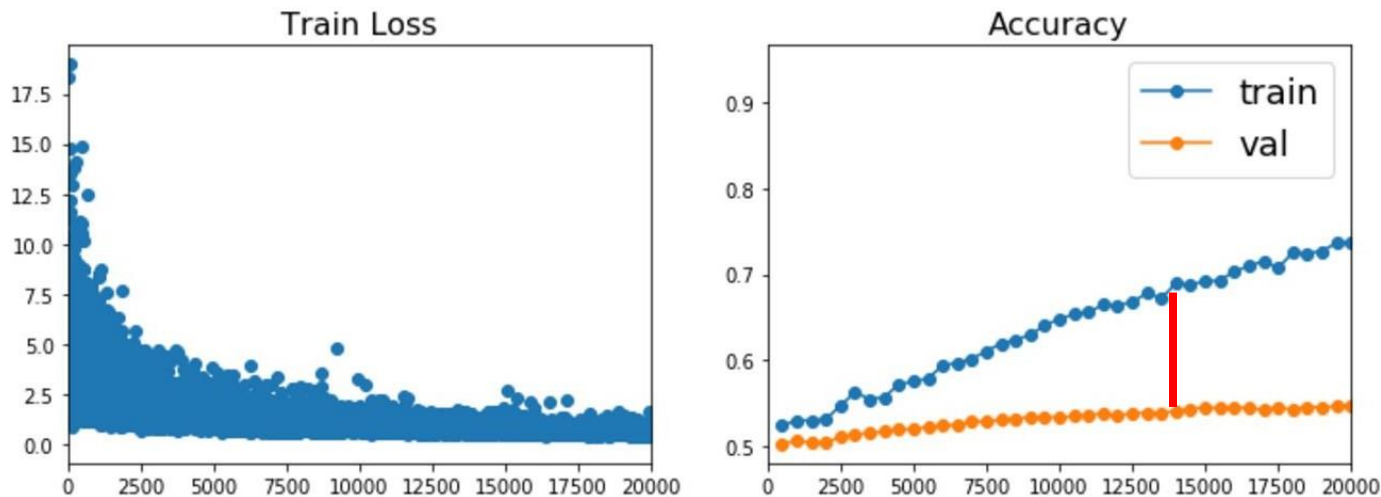
Cyclic learning rate schedules can make this work even better!

Loshchilov and Hutter, "SGDR: Stochastic gradient descent with restarts", arXiv 2016

Huang et al, "Snapshot ensembles: train 1, get M for free", ICLR 2017

Figures copyright Yixuan Li and Geoff Pleiss, 2017. Reproduced with permission.

# How to improve single-model performance?



Regularization

# Regularization

# REGULARIZATION

## Problem of Overfitting:

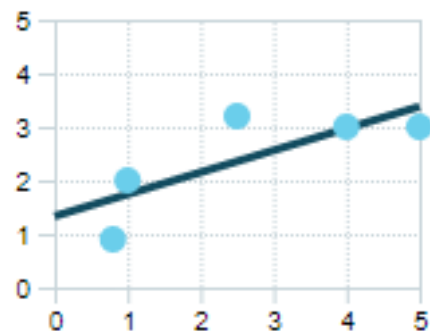
When the model fits the training data but does not have a good predicting performance and generalization power, we have an overfitting problem.

- Regularization is a technique used to avoid this overfitting problem.
- The idea behind regularization is that models that overfit the data are complex models that have for example **too many parameters**.
- In the example below we see how three different models fit the same dataset.
- We used different degrees of polynomials : 1 (linear), 2 (quadratic) and 3 (cubic).
- Notice how the cubic polynomial "sticks" to the data but does not describe the underlying relationship of the data points.

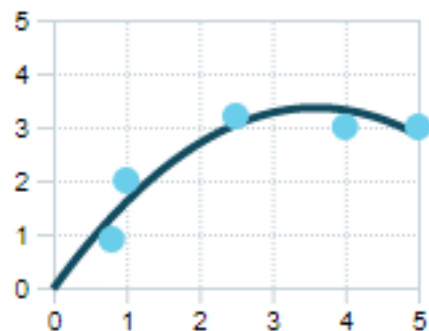


# REGULARIZATION

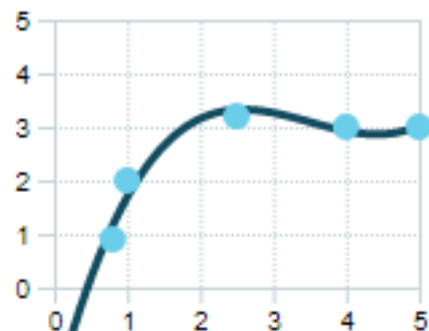
## Different fitting models



Polynomial : Degree 1



Polynomial : Degree 2



Polynomial : Degree 3

# REGULARIZATION

Cost function visualization



But very complex and overfits

# HOW DOES REGULARIZATION WORKS?

The idea is to **penalize the loss function by adding a complexity term that would give a bigger loss for more complex models**. In our case we can use the square sum of the polynomial parameters.

$$L = \sum_m \left( \sum_{i=0} (\alpha_i x_m^i) - y_m \right)^2 + \lambda \sum_{i=1} \alpha_i^2$$

Adjust/Tune the value of Lambda to penalize more or less the complex models.

**Large lambda** → models with high complexity are ruled out.

**small lambda** → models with high training errors are ruled out.

The **optimal solution lies somewhere in the middle**.

# COMMON REGULARIZATION TECHNIQUES

$$L = \frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1) + \lambda R(W)$$

In common use:

**L2 regularization**

$$R(W) = \sum_k \sum_l W_{k,l}^2 \quad (\text{Weight decay})$$

L1 regularization

$$R(W) = \sum_k \sum_l |W_{k,l}|$$

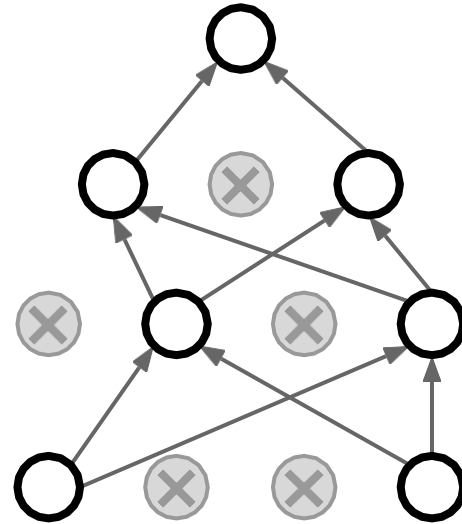
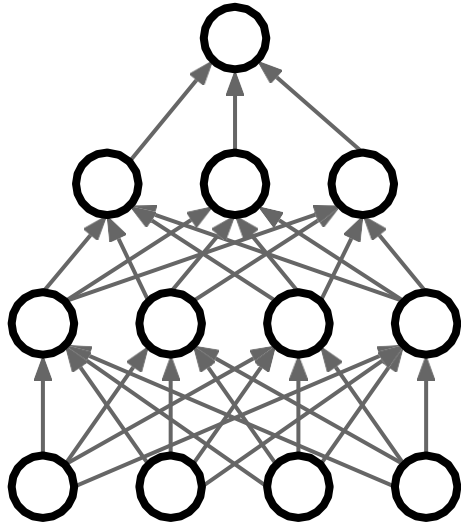
Elastic net (L1 + L2)

$$R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$$

Dropout (Cont..)

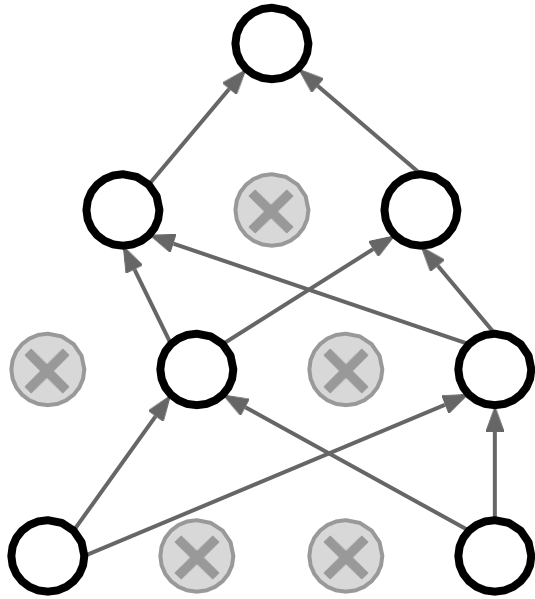
# REGULARIZATION: DROPOUT

IN EACH FORWARD PASS, RANDOMLY SET SOME NEURONS TO ZERO PROBABILITY OF DROPPING IS A HYPERPARAMETER; 0.5 IS COMMON



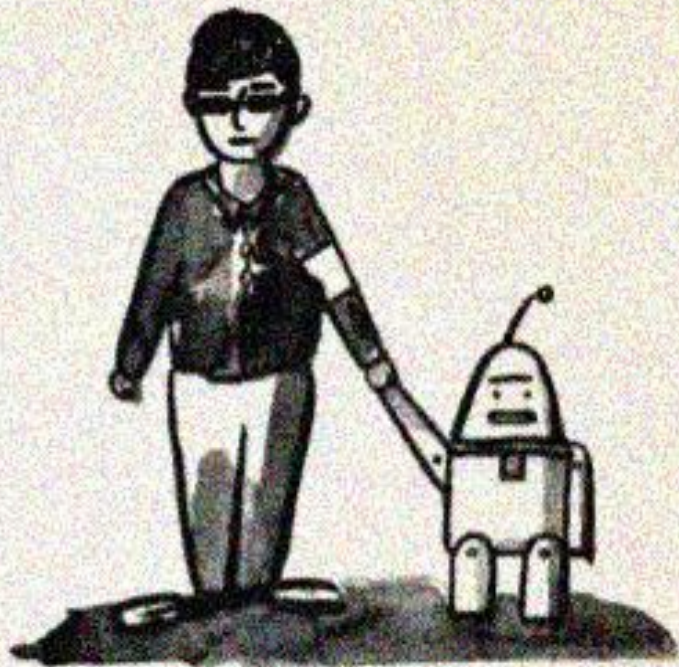
# REGULARIZATION: DROPOUT

How can this possibly be a good idea?



Prevents co-adaptation of features





THANK YOU!

QUESTIONS?